

The Weeding Company

ADDITIONAL QUESTION

What's good about your current design

- **Clear separation:** a *master* DB for global metadata (orgs + admins) and *tenant* collections for each org. This keeps global data centralized and tenant data isolated.
 - **Dynamic tenant creation:** creating `org_<name>` collections programmatically is simple and works for many use cases.
 - **JWT auth:** stateless tokens are easy to scale and work well for APIs.
 - **BCrypt for passwords:** good practice — passwords are hashed, not stored plain.
-

Main trade-offs and limitations (what to watch out for)

1. Single MongoDB instance (single database)

- *Pro:* Easy to implement and operate for small projects.
- *Con:* Single point of failure and potential performance bottleneck as tenants and data grow.

2. One collection per org in the same DB

- *Pro:* Logical separation, easy to query per-tenant.
- *Con:* If you have thousands of orgs, thousands of collections can stress the server and increase metadata overhead. Backup/restore of a single tenant is harder.

3. Auth & RBAC handled manually in controllers

- *Pro:* Simple and explicit.
 - *Con:* Harder to enforce consistently as you add endpoints. Also repetitive code. A centralized JWT filter/middleware is preferable.
-

Practical suggestions to improve

Short-term

1. Centralized JWT filter

- Add a Spring Security filter that validates tokens and attaches claims to request context — controllers read from context rather than parsing the token repeatedly.

2. Create standard template + validation

- Use MongoDB schema validation or a JSON Schema per tenant to avoid garbage data.

Medium-term (scale & safety)

1. Shard or split tenants

- For many tenants or large tenants, consider:
 - grouping multiple tenants per database (DB-per-100-tenants), or
 - using separate databases for very large tenants (db-per-tenant for big customers).

2. Move to replica sets + backup strategy

- Use MongoDB replica sets for high availability and scheduled backups for disaster recovery.

Longer-term (enterprise)

1. Multi-region deployments

- Use geo-aware routing and regional DB replicas for low latency.

2. Policy / billing / audit

- Add audit logs per tenant, quota management, and billing hooks.

Summary

The current design is a good starting point: it keeps global metadata in a master database and creates separate tenant collections dynamically. That makes the code simple and easy to test. For production, I'd add a centralized JWT filter, schema validation, and move MongoDB to a replica set. If we expect thousands

of tenants or very large customers, consider a hybrid model (shared collections for small tenants, dedicated DB or cluster for big tenants). These steps give us better availability, clearer isolation, and easier operational control.
